

BackPropagation

There will be some functions that start with the word "grader" ex: grader_sigmoid(), grader_forwardprop(), grader_backprop() etc, you should not change those function definition.

Every Grader function has to return True.

Loading data ¶

In [1]:

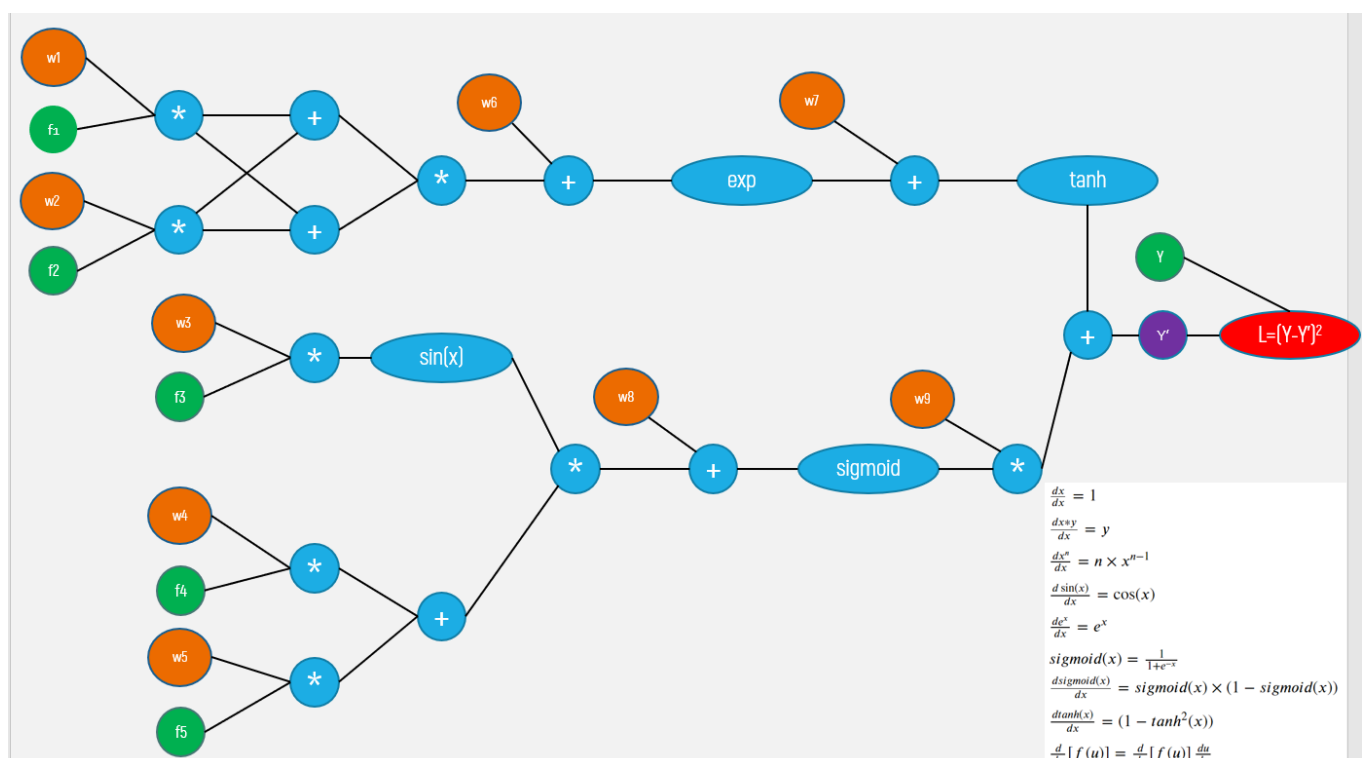
```
import pickle
import numpy as np
import math
from tqdm import tqdm
import matplotlib.pyplot as plt

with open('data.pkl', 'rb') as f:
    data = pickle.load(f)
print(data.shape)
X = data[:, :5]
y = data[:, -1]
print(X.shape, y.shape)
```

(506, 6)

(506, 5) (506,)

Computational graph



- If you observe the graph, we are having input features [f1, f2, f3, f4, f5] and 9 weights [w1, w2, w3, w4, w5, w6, w7, w8, w9].
- The final output of this graph is a value L which is computed as $(Y - Y')^2$

Task 1: Implementing backpropagation and Gradient checking

Check this video for better understanding of the computational graphs and back propagation

In [2]:

```
from IPython.display import YouTubeVideo
YouTubeVideo('i940vYb6noo',width="1000",height="500")
```

Out[2]:

CS231n Winter 2016: Lecture 4: Backpropagation, Neural Network




- Write two functions


- Forward propagation (Write your code in `def forward_propagation()`)

For easy debugging, we will break the computational graph into 3 parts.

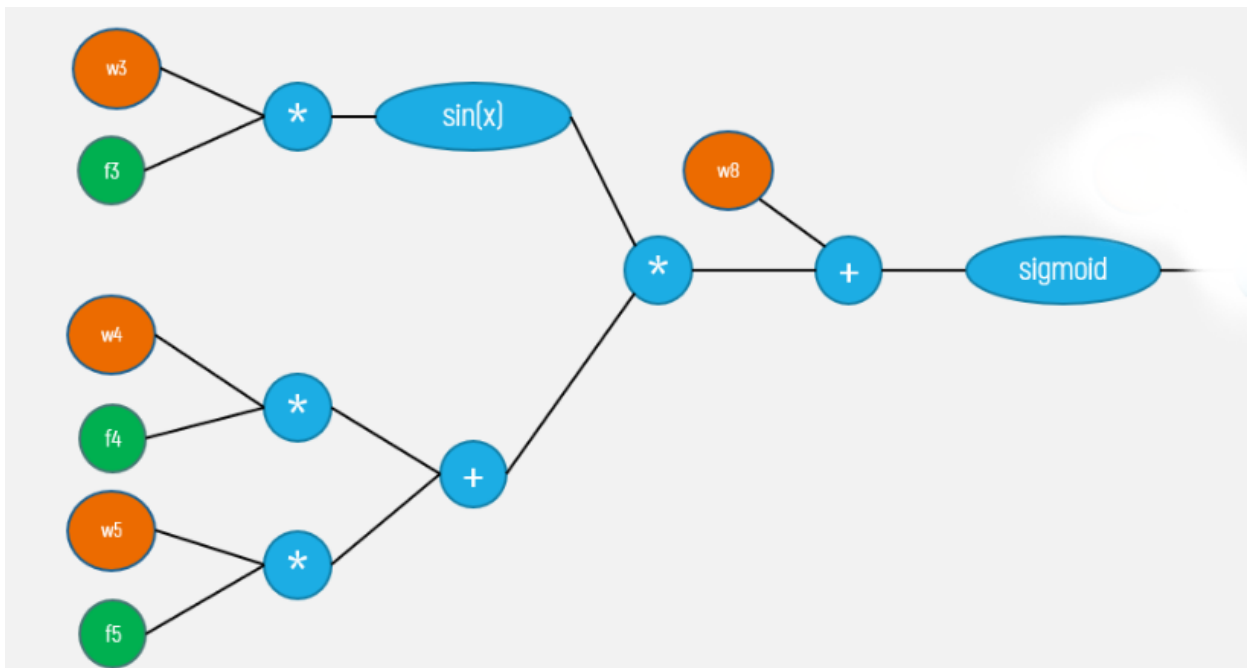
Part 1



Part 2



Part 3



```
def forward_propagation(X, y, W):

    # X: input data point, note that in this assignment you are having 5-d data points
    # y: output variable
    # W: weight array, its of length 9, W[0] corresponds to w1 in graph, W[1] corresponds to w2 in graph,
    # ..., W[8] corresponds to w9 in graph.
    # you have to return the following variables
    # exp= part1 (compute the forward propagation until exp and then store the values in exp)
    # tanh =part2(compute the forward propagation until tanh and then store the values in tanh)
    # sig = part3(compute the forward propagation until sigmoid and then store the values in sig)
    # now compute remaining values from computational graph and get y'
    # write code to compute the value of  $L=(y-y')^2$ 
    # compute derivative of L w.r.to Y' and store it in dl
    # Create a dictionary to store all the intermediate values
    # store L, exp,tanh,sig,dl variables

    return (dictionary, which you might need to use for back propagation)
```

- Backward propagation(Write your code in `def backward_propagation()`)

```
def backward_propagation(L, W,dictionary):

    # L: the loss we calculated for the current point
    # dictionary: the outputs of the forward_propagation() function
    # write code to compute the gradients of each weight [w1,w2,w3,...,w9]
    # Hint: you can use dict type to store the required variables
    # return dW, dW is a dictionary with gradients of all the weights

    return dW
```

Gradient clipping

Check this [blog link \(https://towardsdatascience.com/how-to-debug-a-neural-network-with-gradient-checking-41deec0357a9\)](https://towardsdatascience.com/how-to-debug-a-neural-network-with-gradient-checking-41deec0357a9) for more details on Gradient clipping

we know that the derivative of any function is

$$\lim_{\epsilon \rightarrow 0} \frac{f(x + \epsilon) - f(x - \epsilon)}{2\epsilon}$$

- The definition above can be used as a numerical approximation of the derivative. Taking an epsilon small enough, the calculated approximation will have an error in the range of epsilon squared.
- In other words, if epsilon is 0.001, the approximation will be off by 0.00001.

Therefore, we can use this to approximate the gradient, and in turn make sure that backpropagation is implemented properly. This forms the basis of gradient checking!

Gradient checking example

lets understand the concept with a simple example: $f(w_1, w_2, x_1, x_2) = w_1^2 \cdot x_1 + w_2 \cdot x_2$
 from the above function , lets assume $w_1 = 1, w_2 = 2, x_1 = 3, x_2 = 4$ the gradient of f w.r.t w_1 is

$$\begin{aligned}\frac{df}{dw_1} &= dw_1 = 2 \cdot w_1 \cdot x_1 \\ &= 2 \cdot 1 \cdot 3 \\ &= 6\end{aligned}$$

let calculate the approximate gradient of w_1 as mentioned in the above formula and considering $\epsilon = 0.0001$

$$\begin{aligned}dw_1^{approx} &= \frac{f(w_1+\epsilon, w_2, x_1, x_2) - f(w_1-\epsilon, w_2, x_1, x_2)}{2\epsilon} \\ &= \frac{((1+0.0001)^2 \cdot 3 + 2 \cdot 4) - ((1-0.0001)^2 \cdot 3 + 2 \cdot 4)}{2\epsilon} \\ &= \frac{(1.00020001 \cdot 3 + 2 \cdot 4) - (0.99980001 \cdot 3 + 2 \cdot 4)}{2 \cdot 0.0001} \\ &= \frac{(11.00060003) - (10.99940003)}{0.0002} \\ &= 5.999999999999\end{aligned}$$

Then, we apply the following formula for gradient check: $gradient_check = \frac{\|(dW - dW^{approx})\|_2}{\|(dW)\|_2 + \|(dW^{approx})\|_2}$

The equation above is basically the Euclidean distance normalized by the sum of the norm of the vectors. We use normalization in case that one of the vectors is very small. As a value for epsilon, we usually opt for $1e-7$. Therefore, if gradient check return a value less than $1e-7$, then it means that backpropagation was implemented correctly. Otherwise, there is potentially a mistake in your implementation. If the value exceeds $1e-3$, then you are sure that the code is not correct.

in our example: $gradient_check = \frac{(6 - 5.999999999994898)}{(6 + 5.999999999994898)} = 4.2514140356330737e^{-13}$

you can mathamatically derive the same thing like this

$$\begin{aligned}dw_1^{approx} &= \frac{f(w_1+\epsilon, w_2, x_1, x_2) - f(w_1-\epsilon, w_2, x_1, x_2)}{2\epsilon} \\ &= \frac{((w_1+\epsilon)^2 \cdot x_1 + w_2 \cdot x_2) - ((w_1-\epsilon)^2 \cdot x_1 + w_2 \cdot x_2)}{2\epsilon} \\ &= \frac{4 \cdot \epsilon \cdot w_1 \cdot x_1}{2\epsilon} \\ &= 2 \cdot w_1 \cdot x_1\end{aligned}$$

Implement Gradient checking

(Write your code in `def gradient_checking()`)

Algorithm

```

W = initialize_randomly
def gradient_checking(data_point, W):

    # compute the L value using forward_propagation()
    # compute the gradients of W using backward_propagation()
    approx_gradients = []
    for each wi weight value in W:
        # add a small value to weight wi, and then find the values of L with
        the updated weights
        # subtract a small value to weight wi, and then find the values of L
        with the updated weights
        # compute the approximation gradients of weight wi
        approx_gradients.append(approximation gradients of weight wi)
    # compare the gradient of weights W from backward_propagation() with the
    approximation gradients of weights with gradient_check formula
    return gradient_check

```

NOTE: you can do sanity check by checking all the return values of gradient_checking(),
they have to be zero. if not you have bug in your code

Task 2 : Optimizers

- As a part of this task, you will be implementing 3 type of optimizers(methods to update weight)
- Use the same computational graph that was mentioned above to do this task
- Initilze the 9 weights from normal distribution with mean=0 and std=0.01

Check below video and [this \(https://cs231n.github.io/neural-networks-3/\)](https://cs231n.github.io/neural-networks-3/) blog

In [3]:

```
from IPython.display import YouTubeVideo
YouTubeVideo('gYpoJmIgyXA',width="1000",height="500")
```

Out[3]:

CS231n Winter 2016: Lecture 5: Neural Networks Part 2



Algorithm

```
for each epoch(1-100):
    for each data point in your data:
        using the functions forward_propagation() and backward_propagation()
compute the gradients of weights
    update the weights with help of gradients ex: w1 = w1-learning_rate*dw1
```

Implement below tasks

- **Task 2.1:** you will be implementing the above algorithm with **Vanilla** update of weights
- **Task 2.2:** you will be implementing the above algorithm with **Momentum** update of weights
- **Task 2.3:** you will be implementing the above algorithm with **Adam** update of weights

Note : If you get any assertion error while running grader functions, please print the variables in grader functions and check which variable is returning False .Recheck your logic for that variable .

Task 1

Forward propagation

In [4]:

```

def sigmoid(activation):
    '''In this function, we will compute the sigmoid(z)'''
    # we can use this function in forward and backward propagation

    return 1 / (1 + np.exp(-activation))

def forward_propagation(x, y,Weights):
    '''In this function, we will compute the forward propagation '''
    # X: input data point, note that in this assignment you are having 5-d data points
    # y: output variable
    # W: weight array, its of length 9, W[0] corresponds to w1 in graph, W[1] corresponds to w2 in graph,..., W[8] corresponds to w9 in graph.
    # you have to return the following variables
    # exp= part1 (compute the forward propagation until exp and then store the values in exp)
    # tanh =part2(compute the forward propagation until tanh and then store the values in tanh)
    # sig = part3(compute the forward propagation until sigmoid and then store the values in sig)
    # now compute remaining values from computational graph and get y'
    # write code to compute the value of L=(y-y')^2
    # compute derivative of L w.r.to Y' and store it in dl
    # Create a dictionary to store all the intermediate values
    # store L, exp,tanh,sig variables

    activation1 =(Weights[0] * X[0][0] + Weights[1] * X[0][1]) * (Weights[0] * X[0][0] + Weights[1] * X[0][1])
    activation1 = activation1 + Weights[5]
    Exp = np.exp(activation1)
    #Exp_derv = np.exp(activation1)
    activation2 = Exp + Weights[6]
    Tanh = np.tanh(activation2)
    #Tanh_derv=1-np.tanh(activation2)**2
    activation3 =(Weights[2] * X[0][2])
    sin=math.sin(activation3)
    #sin_derv=np.cos(activation3)
    #activation4 =(Weights[4] * X[0][4] + Weights[5] * X[0][4])
    activation4 =(Weights[3] * X[0][3] + Weights[4] * X[0][4])
    activation5=(activation4 * sin) + Weights[7]
    Sig= sigmoid(activation5)
    #Sig_derv= sigmoid(activation5) * (1-sigmoid(activation5))
    y_=(Sig * Weights[8]) + Tanh
    loss = ((y-y_)**2)
    dl = -2 * (y-y_)
    list_var = ['loss','exp','tanh','sigmoid','dl','sin']
    list_values=[loss,Exp,Tanh,Sig,dl,sin]

    dict_forward_prop = dict(zip(list_var, list_values))
    return (dict(dict_forward_prop))

```

Grader function - 1

In [5]:

```
def grader_sigmoid(z):  
    val=sigmoid(z)  
    assert(val==0.8807970779778823)  
    return True  
grader_sigmoid(2)
```

Out[5]:

True

Grader function - 2

In [6]:

```
def grader_forwardprop(data):  
    d1 = (np.round(data['d1'],4)==-1.9285)  
    loss=(np.round(data['loss'],4)==0.9298)  
    part1=(np.round(data['exp'],4)==1.1273)  
    part2=(np.round(data['tanh'],4)==0.8418)  
    part3=(np.round(data['sigmoid'],4)==0.5279)  
    assert(d1 and loss and part1 and part2 and part3)  
    return True  
w=np.ones(9)*0.1  
d1=forward_propagation(X[0],y[0],w)  
print(d1)  
grader_forwardprop(d1)
```

```
{'loss': 0.9298048963072919, 'exp': 1.1272967040973583, 'tanh': 0.84179341  
92562146, 'sigmoid': 0.5279179387419721, 'd1': -1.9285278284819143, 'sin':  
-0.14538296400984968}
```

Out[6]:

True

Backward propagation

In [7]:

```

def backward_propagation(x,Weights,dict_forward_prop,y):
    '''In this function, we will compute the backward propagation '''
    # L: the loss we calculated for the current point
    # dictionary: the outputs of the forward_propagation() function
    # write code to compute the gradients of each weight [w1,w2,w3,...,w9]
    # Hint: you can use dict type to store the required variables
    # dw1 = # in dw1 compute derivative of L w.r.to w1
    # dw2 = # in dw2 compute derivative of L w.r.to w2
    # dw3 = # in dw3 compute derivative of L w.r.to w3
    # dw4 = # in dw4 compute derivative of L w.r.to w4
    # dw5 = # in dw5 compute derivative of L w.r.to w5
    # dw6 = # in dw6 compute derivative of L w.r.to w6
    # dw7 = # in dw7 compute derivative of L w.r.to w7
    # dw8 = # in dw8 compute derivative of L w.r.to w8
    # dw9 = # in dw9 compute derivative of L w.r.to w9
    p=Weights[0] * X[0][0] + Weights[1] * X[0][1]
    q=Weights[0] * X[0][0] + Weights[1] * X[0][1]
    i = p * q
    h= i + Weights[5]
    a = np.exp(h)
    Exp=a
    Exp_derv = np.exp(h)
    b = a + Weights[6]
    c=np.tanh(b)
    Tanh=c
    Tanh_derv=1-np.tanh(b)**2
    k=(Weights[2] * X[0][2])
    j=np.sin(k)
    Sin=j
    Sin_derv=np.cos(k)
    n =(Weights[4] * X[0][3])
    o =(Weights[5] * X[0][4])
    m=n+o
    g=j*m
    f=g + Weights[7]
    d= sigmoid(f)
    sig=d
    Sig_derv= sigmoid(f) * (1-sigmoid(f))
    e=d*Weights[8]
    y_=e+c
    d1 = -1*(2 * (y-y_))
    dw1=float(d1*Tanh_derv*Exp_derv*(q*X[0][0]+p*X[0][0]))
    dw2=float(d1*Tanh_derv*Exp_derv*(q*X[0][1]+p*X[0][1]))
    dw3=float(d1*Weights[8]*Sig_derv*m*Sin_derv*X[0][2])
    dw4=float(d1*Weights[8]*Sig_derv*j*X[0][3])
    dw5=float(d1*Weights[8]*Sig_derv*j*X[0][4])
    dw6=float(d1*Tanh_derv*Exp_derv)
    dw7=float(d1*Tanh_derv)
    dw8=float(d1*Weights[8]*Sig_derv)
    dw9=float(d1*d)
    list_values1=[]
    list_var1 = ['dw1', 'dw2', 'dw3', 'dw4', 'dw5', 'dw6', 'dw7', 'dw8', 'dw9']
    list_values1=[dw1,dw2,dw3,dw4,dw5,dw6,dw7,dw8,dw9]
    dict_back_prop = dict(zip(list_var1, list_values1))
    return dict(dict_back_prop)

```

Grader function - 3

In [8]:

```
def grader_backprop(data):  
    dw1=(np.round(data['dw1'],4)==-0.2297)  
    dw2=(np.round(data['dw2'],4)==-0.0214)  
    dw3=(np.round(data['dw3'],4)==-0.0056)  
    dw4=(np.round(data['dw4'],4)==-0.0047)  
    dw5=(np.round(data['dw5'],4)==-0.001)  
    dw6=(np.round(data['dw6'],4)==-0.6335)  
    dw7=(np.round(data['dw7'],4)==-0.5619)  
    dw8=(np.round(data['dw8'],4)==-0.0481)  
    dw9=(np.round(data['dw9'],4)==-1.0181)  
    assert(dw1 and dw2 and dw3 and dw4 and dw5 and dw6 and dw7 and dw8 and dw9)  
    return True  
w=np.ones(9)*0.1  
d1=forward_propagation(X[0],y[0],w)  
d1=backward_propagation(X[0],w,d1,y[0])  
grader_backprop(d1)
```

Out[8]:

True

Implement gradient checking

In [9]:

```

def gradient_checking(x,y,W1):
    # compute the L value using forward_propagation()
    # compute the gradients of W using backward_propagation()
    D_forward=forward_propagation(x,y,W1)
    D_back_prop_grad1 = backward_propagation(x,W1,D_forward,y)
    approx_gradients = []
    grad_check_list=[]
    for i,wi in enumerate(W1):
        grad_check=0
        Weights_forward1=W1
        e=0.0001
        approx_grad=0
        Weights_forward1[i] = wi + e
        D_forward=forward_propagation(x,y,Weights_forward1)
        L_plus =D_forward['loss']
        Weights_forward1=W1
        Weights_forward1[i] = wi -e
        D_forward=forward_propagation(x,y,Weights_forward1)
        L_minus =D_forward['loss']
        approx_grad=float((L_plus - L_minus)/ (2*e))
        approx_gradients.append(approx_grad)
        W1[i]=wi
    # compare the gradient of weights W from backward_propagation() with the approximation
    # compare the gradient of weights with gradient_check formula
    values_back_prop = D_back_prop_grad1.values()
    values_back_prop_list = list(values_back_prop)
    for i in range(len(approx_gradients)):
        numerator = np.linalg.norm(values_back_prop_list[i] - approx_gradients[i])
        denominator = np.linalg.norm(values_back_prop_list[i]) + np.linalg.norm(approx_gradients[i])
        difference = numerator / denominator
        if difference < 1e-7:
            print("{} The gradient is correct!".format(difference))
        else:
            print("The gradient is wrong!".format(difference))
    return
W1=np.ones(9)*0.1
x=X[:1]
y_1=y[:1]
gradient_checking(x,y_1,W1)

```

```

1.0370728946337426e-08 The gradient is correct!
8.174982888388155e-11 The gradient is correct!
1.7287700041112022e-09 The gradient is correct!
1.87486944153289e-12 The gradient is correct!
4.2849738752544037e-10 The gradient is correct!
7.610196933782967e-10 The gradient is correct!
3.1480030084674753e-09 The gradient is correct!
4.0368014625577295e-10 The gradient is correct!
3.361951351774315e-13 The gradient is correct!

```

Task 2: Optimizers

Algorithm with Vanilla update of weights

In [10]:

```
def vanilla_update(x, y, W, epoc, N, alpha):
    loss_lst=[]
    for i in tqdm(range(epoc)):
        loss=0
        for step in range(N):
            # Update weights with gradient, scores
            D_forward=forward_propagation(x[step],y[step],W)
            D_back_prop_grad = backward_propagation(x[step],W,D_forward,y[step])
            W[0]=W[0] - alpha * D_back_prop_grad['dw1']
            W[1]=W[1] - alpha * D_back_prop_grad['dw2']
            W[2]=W[2] - alpha * D_back_prop_grad['dw3']
            W[3]=W[4] - alpha * D_back_prop_grad['dw4']
            W[4]=W[5] - alpha * D_back_prop_grad['dw5']
            W[5]=W[5] - alpha * D_back_prop_grad['dw6']
            W[6]=W[6] - alpha * D_back_prop_grad['dw7']
            W[7]=W[7] - alpha * D_back_prop_grad['dw8']
            W[8]=W[8] - alpha * D_back_prop_grad['dw9']
            loss=loss+D_forward['loss']
            loss_avg=loss/len(x)
        loss_lst.append(loss_avg)
    return loss_lst
```

Plot between epochs and loss

```
from tqdm import tqdm
epoc=100
N=len(X)
W = np.random.normal(0.0, 1.0, 9)

loss_lst = vanilla_update(X, y,W,epoc,N, 0.0001)
epoc=np.arange(1,101)
plt.plot(epoc,loss_lst,label='vanilla_update_loss')
plt.legend()
plt.xlabel("epoc: epoc")
plt.ylabel("loss")
plt.title("epoc Vs loss plot")
plt.grid()
plt.show()
```

The plot shows the training loss for a vanilla update method. The x-axis represents the epoch number from 0 to 100, and the y-axis represents the loss from 0.00 to 1.75. The loss starts at approximately 1.85 at epoch 0 and decreases rapidly, reaching a plateau of about 0.05 after epoch 40.

epoc: epoc	vanilla_update_loss
0	1.85
10	0.75
20	0.25
30	0.10
40	0.05
50	0.05
60	0.05
70	0.05
80	0.05
90	0.05
100	0.05

file:///C:/Users/Baskaran Thulukanam/Desktop/New folder/Back_prop - Copy/baskar.mailbox@gmail.com_19_Backpropagation_assignment.ip... 16/21

In [12]:

```

def momentum_update(x, y,W,epoc,N,alpha):
    loss_lst1=[]
    for i in tqdm(range(epoc)):
        loss1=0
        v1,v2,v3,v4,v5,v6,v7,v8,v9=0,0,0,0,0,0,0,0,0
        mu=0.9
        for step in range(N):
            # Update weights with gradient,scores
            D_forward=forward_propagation(x[step],y[step],W)
            D_back_prop_grad = backward_propagation(x[step],W,D_forward,y[step])
            # Momentum update
            v1 = mu * v1 - alpha * D_back_prop_grad[ 'dw1' ]
            W[0] += v1
            v2 = mu * v2 - alpha * D_back_prop_grad[ 'dw2' ]
            W[1] += v2
            v3 = mu * v3 - alpha * D_back_prop_grad[ 'dw3' ]
            W[2] += v3
            v4 = mu * v4 - alpha * D_back_prop_grad[ 'dw4' ]
            W[3] += v4
            v5 = mu * v5 - alpha * D_back_prop_grad[ 'dw5' ]
            W[4] += v5
            v6 = mu * v6 - alpha * D_back_prop_grad[ 'dw6' ]
            W[5] += v6
            v7 = mu * v7 - alpha * D_back_prop_grad[ 'dw7' ]
            W[6] += v7
            v8 = mu * v8 - alpha * D_back_prop_grad[ 'dw8' ]
            W[7] += v8
            v9 = mu * v9 - alpha * D_back_prop_grad[ 'dw9' ]
            W[8] += v9
            loss1=loss1+D_forward[ 'loss' ]
            loss_avg1=loss1/len(x)
        loss_lst1.append(loss_avg1)
    return loss_lst1

```

Plot between epochs and loss

```
from tqdm import tqdm
epoc=100
N=len(X)
W = np.random.normal(0.0, 1.0, 9)

loss_lst1 = momentum_update(X, y,W,epoc,N, 0.0001)
epoc=np.arange(1,101)
plt.plot(epoc,loss_lst1,label='momentum_update_loss')
plt.legend()
plt.xlabel("epoc: epoc")
plt.ylabel("loss")
plt.title("epoc Vs loss plot")
plt.grid()
plt.show()
```

The plot shows the training loss for the momentum update over 100 epochs. The y-axis represents the loss, ranging from 0 to 0.6. The x-axis represents the epoch number, ranging from 0 to 100. The loss starts at approximately 0.6 at epoch 0 and decreases rapidly, reaching a minimum of about 0.05 by epoch 10. After epoch 10, the loss remains relatively constant, fluctuating slightly around 0.05.

epoch: epoc	momentum_update_loss
0	0.60
2	0.20
4	0.08
6	0.05
10	0.05
20	0.05
40	0.05
60	0.05
80	0.05
100	0.05

file:///C:/Users/Baskaran Thulukanam/Desktop/New folder/Back_prop - Copy/baskar.mailbox@gmail.com_19_Backpropagation_assignment.ip... 18/21

In [14]:

```

def adam_update(x, y, W, epoc, N, alpha):
    loss_lst2=[]
    for i in tqdm(range(epoc)):
        loss2=0
        v1,v2,v3,v4,v5,v6,v7,v8,v9=0,0,0,0,0,0,0,0,0
        eps = 1e-8
        beta1 = 0.9
        beta2 = 0.999
        m1,m2,m3,m4,m5,m6,m7,m8,m9=0,0,0,0,0,0,0,0,0
        for step in range(N):
            # Update weights with gradient,scores
            D_forward=forward_propagation(x[step],y[step],W)
            D_back_prop_grad = backward_propagation(x[step],W,D_forward,y[step])
            # adam update
            m1 = beta1*m1 + (1-beta1)* D_back_prop_grad['dw1']
            v1 = beta2*v1 + (1-beta2)*(D_back_prop_grad['dw1'] **2)
            W[0] += - alpha * m1 / (np.sqrt(v1) + eps)
            m2 = beta1*m2 + (1-beta1)* D_back_prop_grad['dw2']
            v2 = beta2*v2 + (1-beta2)*(D_back_prop_grad['dw2'] **2)
            W[1] += - alpha * m2 / (np.sqrt(v2) + eps)
            m3 = beta1*m3 + (1-beta1)* D_back_prop_grad['dw3']
            v3 = beta2*v3 + (1-beta2)*(D_back_prop_grad['dw3'] **2)
            W[2] += - alpha * m3 / (np.sqrt(v3) + eps)
            m4 = beta1*m4 + (1-beta1)* D_back_prop_grad['dw4']
            v4 = beta2*v4 + (1-beta2)*(D_back_prop_grad['dw4'] **2)
            W[3] += - alpha * m4 / (np.sqrt(v4) + eps)
            m5 = beta1*m5 + (1-beta1)* D_back_prop_grad['dw5']
            v5 = beta2*v5 + (1-beta2)*(D_back_prop_grad['dw5'] **2)
            W[4] += - alpha * m5 / (np.sqrt(v5) + eps)
            m6 = beta1*m6 + (1-beta1)* D_back_prop_grad['dw6']
            v6 = beta2*v6 + (1-beta2)*(D_back_prop_grad['dw6'] **2)
            W[5] += - alpha * m6 / (np.sqrt(v6) + eps)
            m7 = beta1*m7 + (1-beta1)* D_back_prop_grad['dw7']
            v7 = beta2*v7 + (1-beta2)*(D_back_prop_grad['dw7'] **2)
            W[6] += - alpha * m7 / (np.sqrt(v7) + eps)
            m8 = beta1*m8 + (1-beta1)* D_back_prop_grad['dw8']
            v8 = beta2*v8 + (1-beta2)*(D_back_prop_grad['dw8'] **2)
            W[7] += - alpha * m8 / (np.sqrt(v8) + eps)
            m9 = beta1*m9 + (1-beta1)* D_back_prop_grad['dw9']
            v9 = beta2*v9 + (1-beta2)*(D_back_prop_grad['dw9'] **2)
            W[8] += - alpha * m9 / (np.sqrt(v9) + eps)

            loss2=loss2+D_forward['loss']
            loss_avg2=loss2/len(x)
            loss_lst2.append(loss_avg2)
    return loss_lst2

```

Plot between epochs and loss

```
from tqdm import tqdm
epoc=100
N=len(X)
W = np.random.normal(0.0, 1.0, 9)

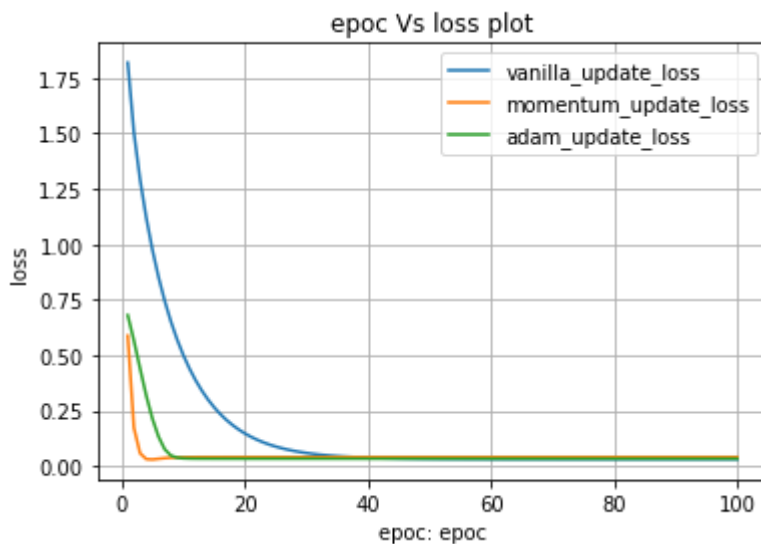
loss_lst2 = adam_update(X, y,W,epoc,N, 0.0001)
epoc=np.arange(1,101)
plt.plot(epoc,loss_lst2,label='adam_update_loss')
plt.legend()
plt.xlabel("epoc: epoc")
plt.ylabel("loss")
plt.title("epoc Vs loss plot")
plt.grid()
plt.show()
```

The plot shows the training loss over 100 epochs. The y-axis, labeled 'loss', ranges from 0.0 to 0.7. The x-axis, labeled 'epoc: epoc', ranges from 0 to 100. A single blue line represents the 'adam_update_loss'. The loss begins at approximately 0.68 at epoch 0, decreases rapidly to about 0.05 by epoch 10, and then remains constant at that level for the rest of the training process.

epoc: epoc	adam_update_loss
0	0.68
5	0.25
10	0.05
20	0.05
40	0.05
60	0.05
80	0.05
100	0.05

In [16]:

```
epoc=np.arange(1,101)
plt.plot(epoc,loss_lst,label='vanilla_update_loss')
plt.plot(epoc,loss_lst1,label='momentum_update_loss')
plt.plot(epoc,loss_lst2,label='adam_update_loss')
plt.legend()
plt.xlabel("epoc: epoc")
plt.ylabel("loss")
plt.title("epoc Vs loss plot")
plt.grid()
plt.show()
```



In []:

In []: